Proposed Solution

Task 1: Document Classification System - Use Case

Context

A bank seeks to automate the classification of customer-submitted PDF documents, like statements, contracts, and applications, by detecting their type and storing them correctly.

System Requirements

Input (1.1)

Documents arrive via:

- Email.
- Web portal.
- Mobile app.
- Third-party API.

Process (1.2)

- **Extract content**: Read text and analyze images.
- **Classify**: Use AI to identify document type.
- **Store**: Save in the right repository with metadata.

Output (1.3)

- Tagged by type (e.g., Contract, Statement)
- Indexed for search
- Logged for compliance

C4 Model Architecture

This solution uses the **C4 Model** with multiple layers to explain all system interactions:

1. Context Diagram (Level 1)

Shows the system in its environment with external actors and systems.

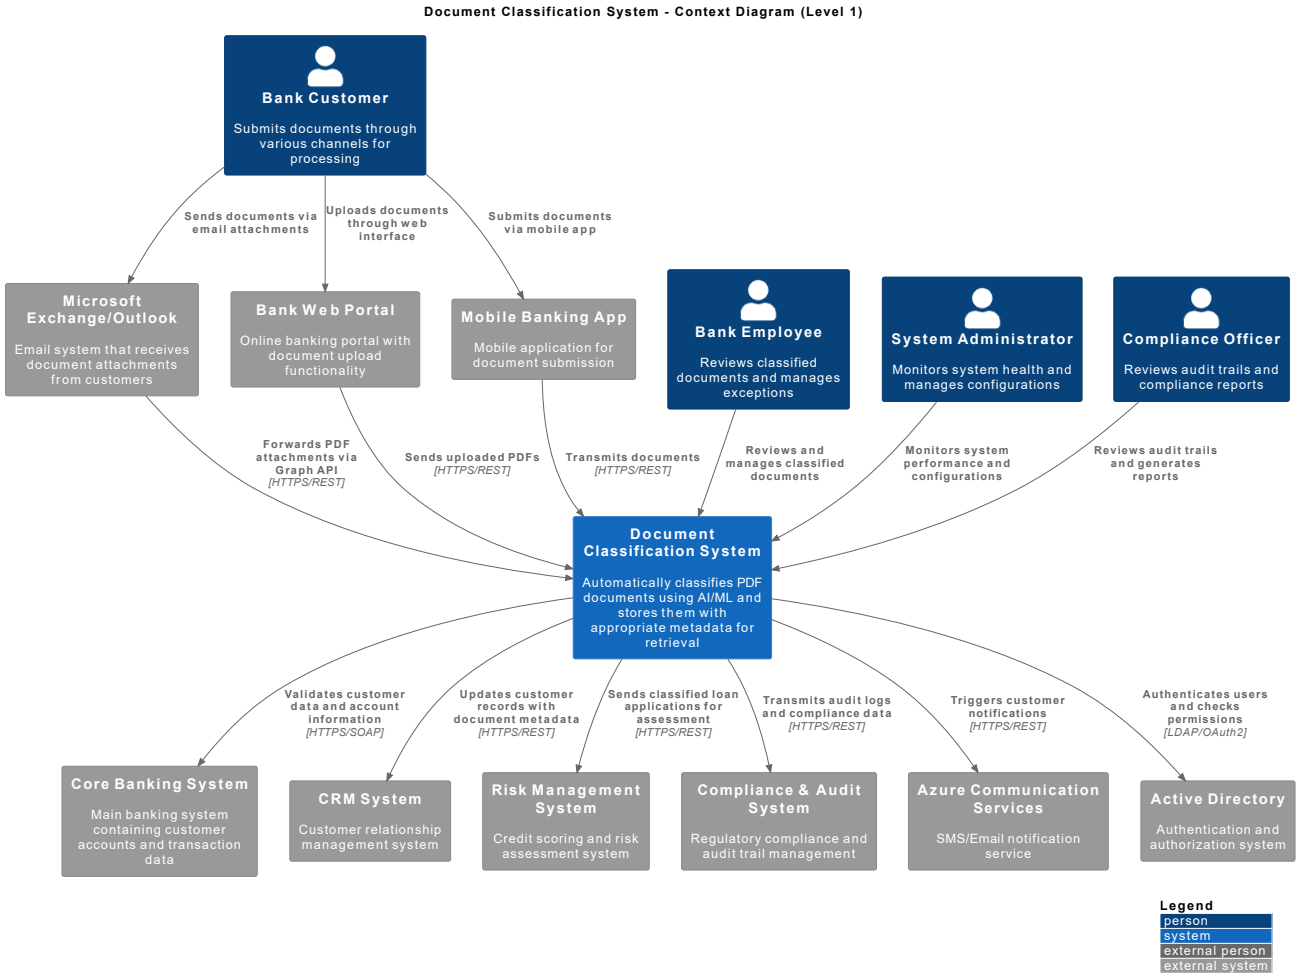

2. Container Diagram (Level 2)

Shows high-level technical building blocks and their interactions.

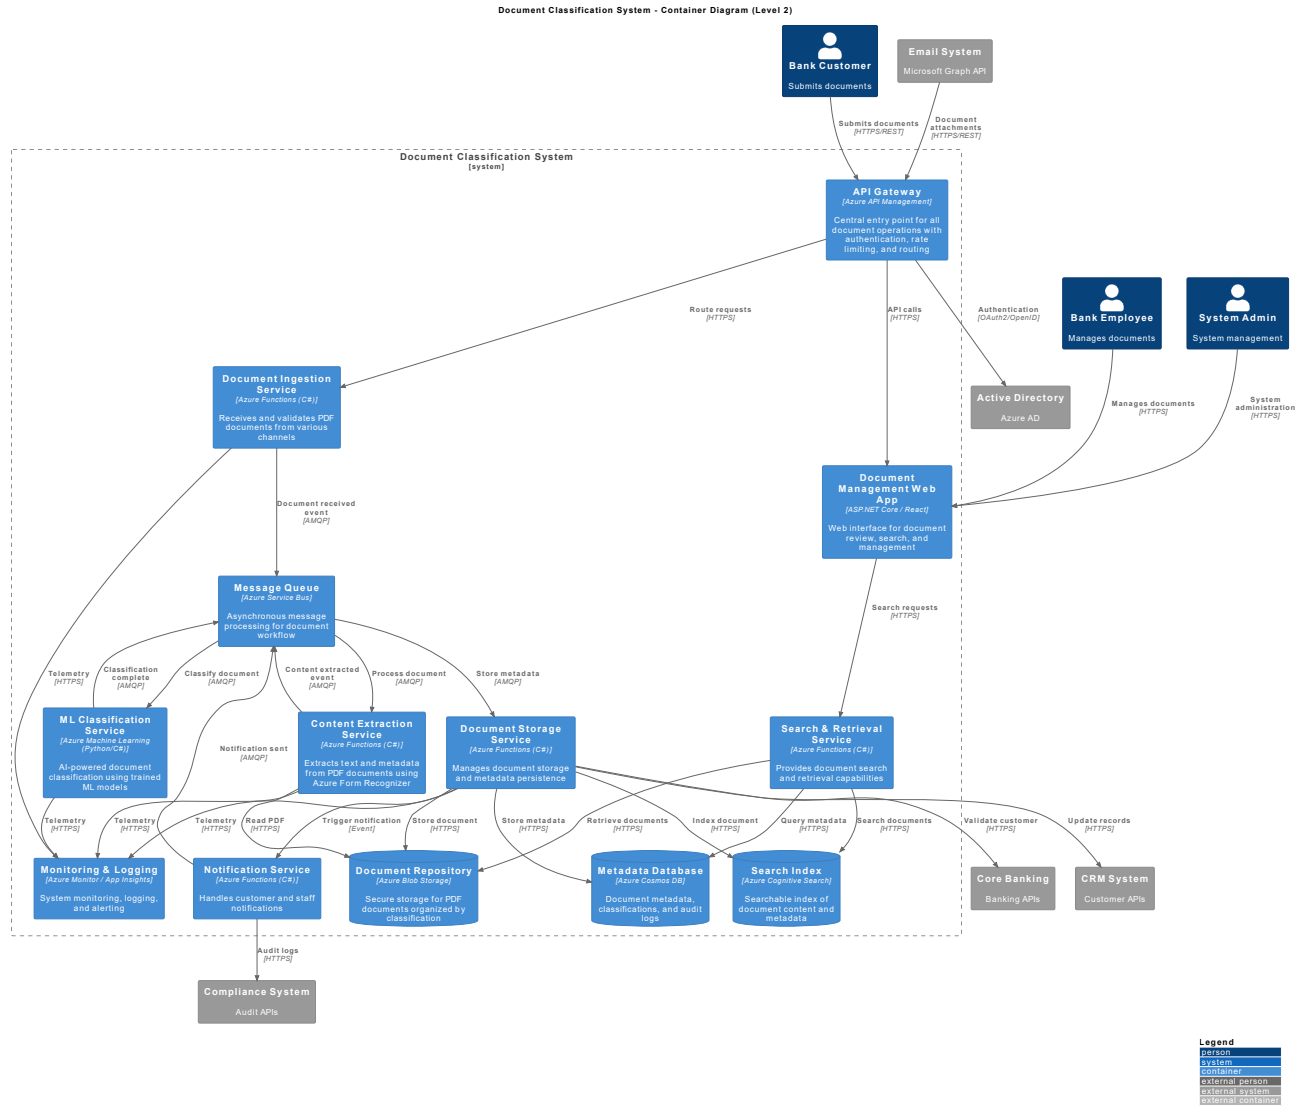

3. Component Diagram (Level 3)

Shows internal components within key containers.

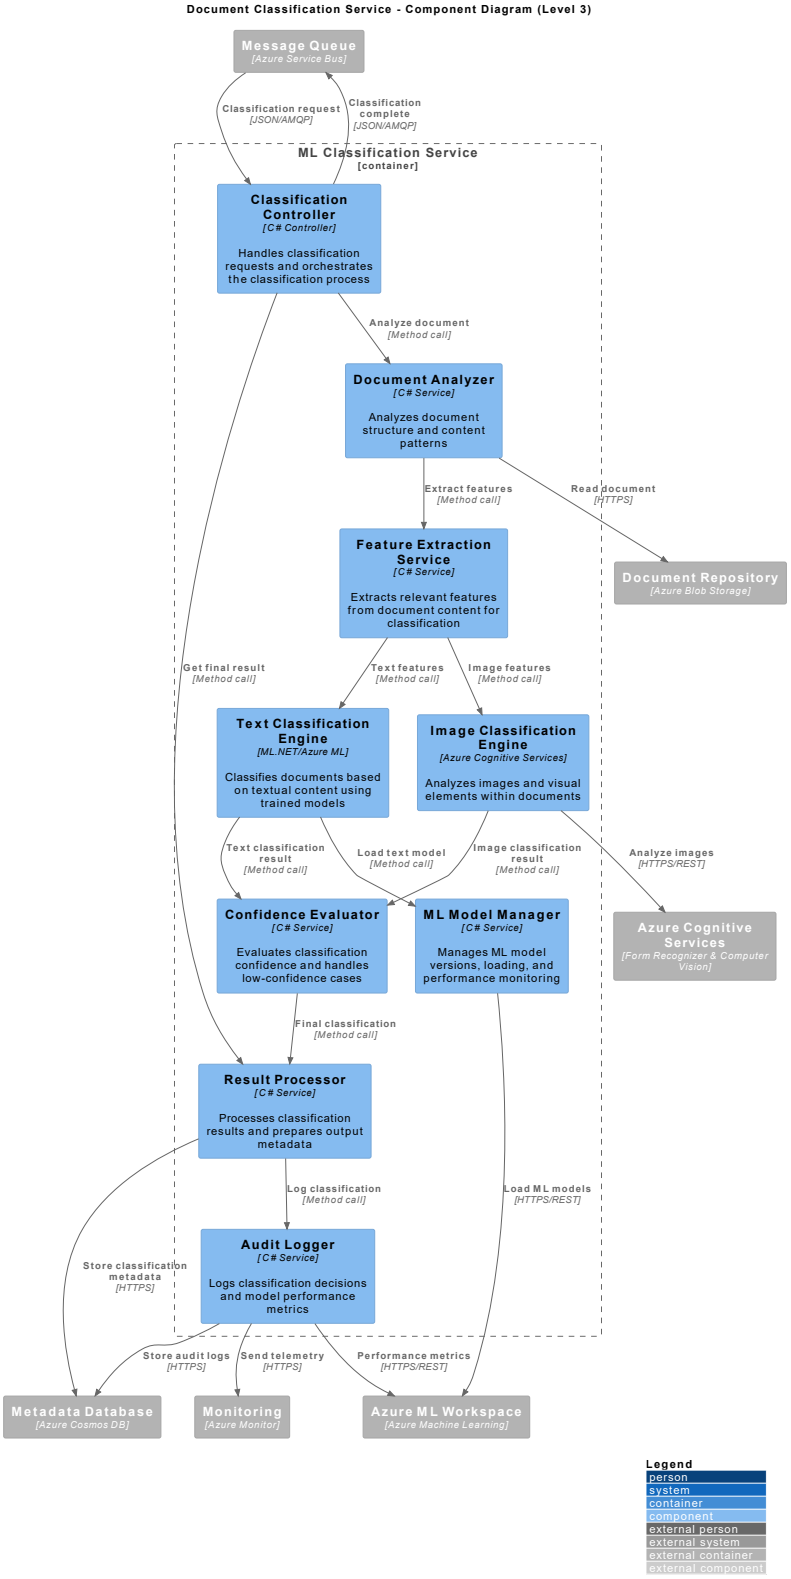

4. Deployment Diagram

Shows infrastructure and deployment architecture on Microsoft Azure.

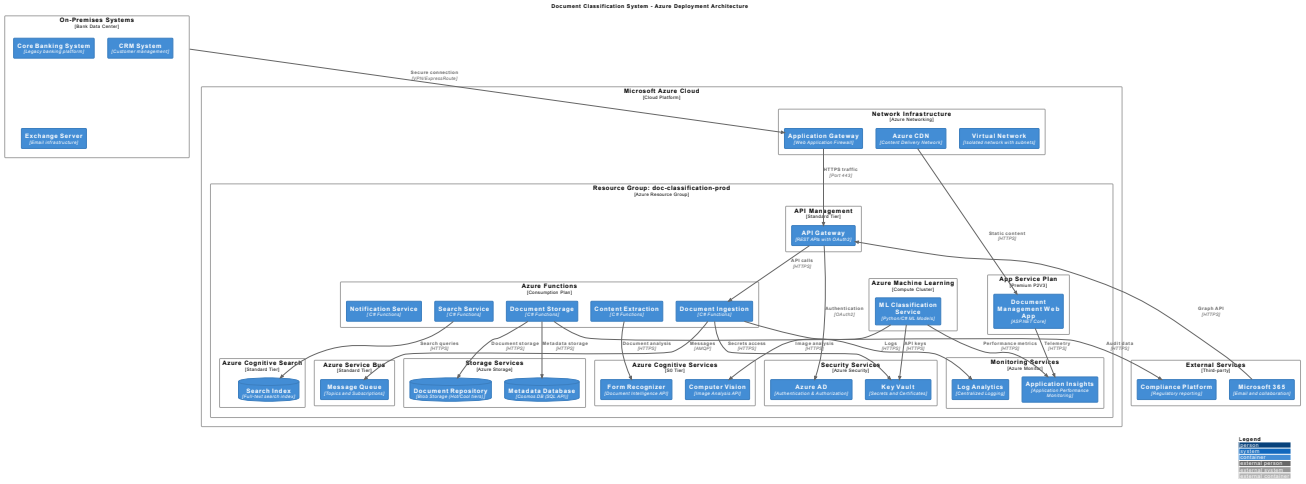

Key Components & Integrations

Core System Components

The solution is built from modular services that work together to classify and store documents efficiently:

- **API Gateway**: Entry point for document submissions.
- **Ingestion Service**: Receives and pre-processes PDFs.
- **ML Engine**: Uses AI to classify documents.
- **Document Repository**: Secure archive for storing files.
- **Metadata Manager**: Handles document tags and indexing.
- **Search Service**: Enables quick and accurate retrieval.

External Integrations

To connect with key systems and ensure smooth operations, the platform integrates with:

- **Email API**: Microsoft Graph for processing attachments.
- **Core Banking API**: Syncs with internal bank systems.
- **CRM API**: Connects customer data for context.
- **Compliance API**: Supports audits and reporting.
- **Notifications**: Azure Communication for user alerts.
- **Cognitive Services**: Powers AI document analysis.

Azure Technology Stack

The solution is cloud-native, built on Microsoft Azure components:

- API gateway via **Azure API Management**.
- Scalable compute with **Azure Functions**.
- AI-powered insights from **Cognitive Services**.
- Custom ML models with **Azure Machine Learning**.
- Secure storage in **Azure Blob Storage**.
- Metadata housed in **Azure Cosmos DB**.
- Search powered by **Azure Search**.
- Messaging via **Azure Service Bus**.
- Secrets managed with **Azure Key Vault**.
- Real-time monitoring using **Azure Monitor**.

Task 2: FinHub Data Processing Solution - From HLD to LLD

Overview

FinHub needs a data platform to collect financial data from multiple APIs (BCRA, ARCA, etc.), run ETL workflows, and deliver analytics to business leaders.

2.1 High-Level Design (HLD)

Architecture Overview

1. Ingestion

- Central API Gateway.
- Connectors for BCRA, ARCA, market data.
- Scheduling service and data validation.

2. Processing

- ETL Orchestrator.
- Data cleaning and enrichment.
- Business rule application.
- Quality checks and alerts.

3. Storage

- Data Lake for raw info.
- Data Warehouse for processed results.
- Metadata and archive layers.

4. Analytics & Reports

- Analytics Engine.
- Scheduled & on-demand reports.
- Live dashboards.
- Exposed via REST APIs.

5. Security & Compliance

- User authentication.
- Audit logs.
- Encryption for data at rest and in transit.
- GDPR/SOX compliance monitoring.

6. Operations

- Monitoring and alerts.
- Central config.

- Backup and recovery.

Data Sources & Flow

Sources

- Central banks, tax agencies, market feeds.
- Internal systems and credit bureaus.

Flow

Ingest → Validate → Transform → Store → Analyze → Report.

Design Considerations

Scalability

- Microservices and auto-scaling.
- Redis caching and load balancing.

Security

- Zero Trust model.
- API protection via OAuth2.
- Secure networking and secret storage.

Compliance

- Governance, retention policies.
- Full audit trails.
- GDPR data privacy tools.

2.2 Low-Level Design (LLD)

Selected Components - Detailed View

Data Ingestion Service

Designed with modularity and flexibility in mind, this service uses factory and strategy patterns to handle multiple data sources.

Main Classes

- `DataIngestionController` : Receives API requests.
- `DataOrchestrator` : Oversees the ingestion flow.
- `DataConnectorFactory` : Instantiates connectors based on source type.
- Specialized connectors: `BCRAConnector` , `ARCAConnector` , `MarketDataConnector` .
- `DataValidator` : Checks data integrity and applies rules.
- `AzureStorageService` : Manages cloud storage.
- `ServiceBusMessageQueue` : Publishes events asynchronously.

Design Patterns Used

- Factory: For connector creation.
- Strategy: Custom validation per source.
- Dependency Injection: Improves modularity and testability.
- Repository: Abstracts data access.

Class Interactions

- Controller depends on `Orchestrator` and `Validator` .
- Orchestrator manages flow using `Factory` , `Validator` , `Storage` , `Queue`
- Connectors implement a shared `IDataConnector` interface.

Data Processing Engine

A robust ETL pipeline, built to support transformation, rule application, and quality control.

Main Components

- `ETLOrchestrator` : Manages full ETL lifecycle.
- `TransformationService` : Performs data cleanup and enhancement.
- `TransformationEngine` : Executes transformation rules
- `BusinessRulesEngine` : Applies financial logic.
- `FinancialCalculationService` : Handles calculations.
- `DataQualityMonitor` : Flags anomalies.
- `SynapseDataWarehouseService` : Stores results.
- `ProcessingJobManager` : Manages ETL jobs.

Patterns Applied

- Chain of Responsibility: Rule processing flow.
- Strategy: Dynamic rule categories.
- Observer: Real-time quality alerts.
- Command: Task execution and tracking.
- Template Method: Structured ETL steps.

Workflow Summary

1. Ingestion request hits `ETLOrchestrator`.
2. Data gets validated.
3. Transformation and enrichment are applied.
4. Business rules and calculations executed.
5. Quality checks run.
6. Processed data stored with full lineage.

2.3 Architectural Enabler Features

1. Azure Data Lake Storage Gen2 Integration

Purpose
To handle large volumes of financial data with scalable, secure storage optimized for analytics and regulatory compliance.

Design Highlights

- **Hierarchical Namespace**: Faster directory-level operations for analytics.
- **Multi-Protocol**: Compatible with Blob APIs and HDFS.
- **Tiered Storage**: Lifecycle automation across Hot, Cool, Archive.
- **Security**: Azure AD integration, encryption, private endpoints.
- **Performance**: Supports parallel access and large file handling.

Implementation Phases

1. Partitioned setup.
2. Lifecycle and tiering automation.
3. Security and compliance hardening.
4. Performance tuning and monitoring.

Key Features

- Auto-partitioning by date/source.
- Smart tiering based on access patterns.
- Synapse integration for analytics.
- Audit logs and granular access control.

2. Azure Service Bus Message Queue

Purpose
To ensure reliable, ordered delivery of financial messages across decoupled services.

Design Highlights

- **Topics/Subs**: Multi-consumer message routing.
- **Message Sessions**: Maintain transaction order.
- **Dead Letter Queues**: Retry failed messages.
- **Duplicate Detection**: Ensure data consistency.
- **Auto-Scaling**: Adjusts to queue load dynamically.

Implementation Phases

1. Core messaging setup.
2. Advanced features (DLQ, sessions, deduplication).
3. Monitoring & alerting.

4. Performance and scaling optimization.

Key Features

- At-least-once delivery.
- Retry logic with exponential backoff.
- Correlated messaging for data tracing.
- Full observability via Azure Monitor.

3. Monitoring with Azure Monitor

Purpose
To ensure visibility into system health, performance, and data quality, with alerts and automation for proactive issue resolution.

Design Highlights

- **Multi-Layer Monitoring**: Infra, app, business metrics.
- **Real-Time Alerts**: Issues surfaced immediately.
- **Distributed Tracing**: Full visibility across services.
- **Custom Dashboards**: KPIs tailored to business needs.
- **Automation**: Enables self-healing workflows.

Implementation Phases

1. Basic system monitoring.
2. Business metrics and quality tracking.
3. Predictive alerts and analytics.
4. Automated remediation.

Key Features

- Telemetry for data pipelines.
- Anomaly detection on quality signals.
- KPI dashboards for executives.
- Integration with incident management tools.

Diagrams

High-Level Design architecture diagram

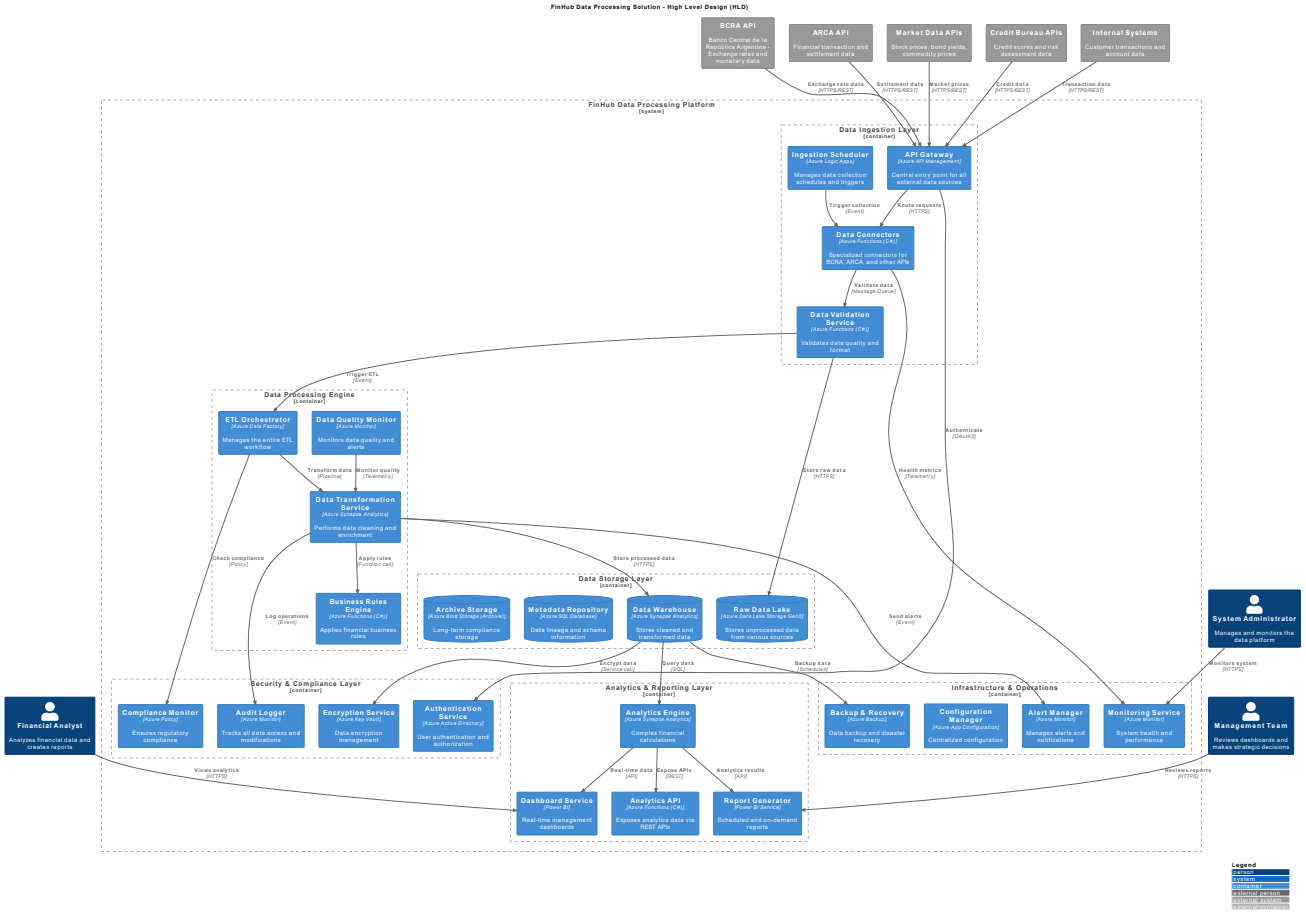

Data Ingestion Service class diagram

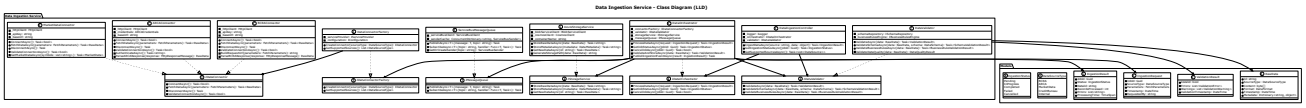

Data Processing Engine class diagram

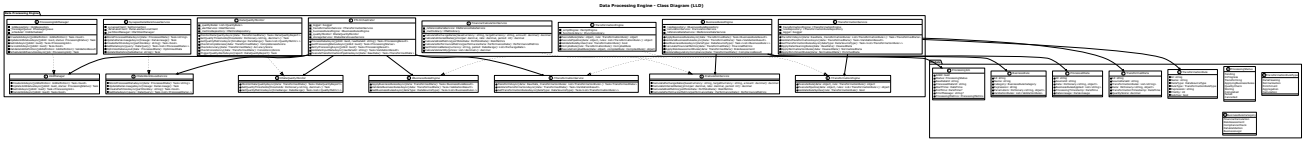

Data ingestion sequence diagram

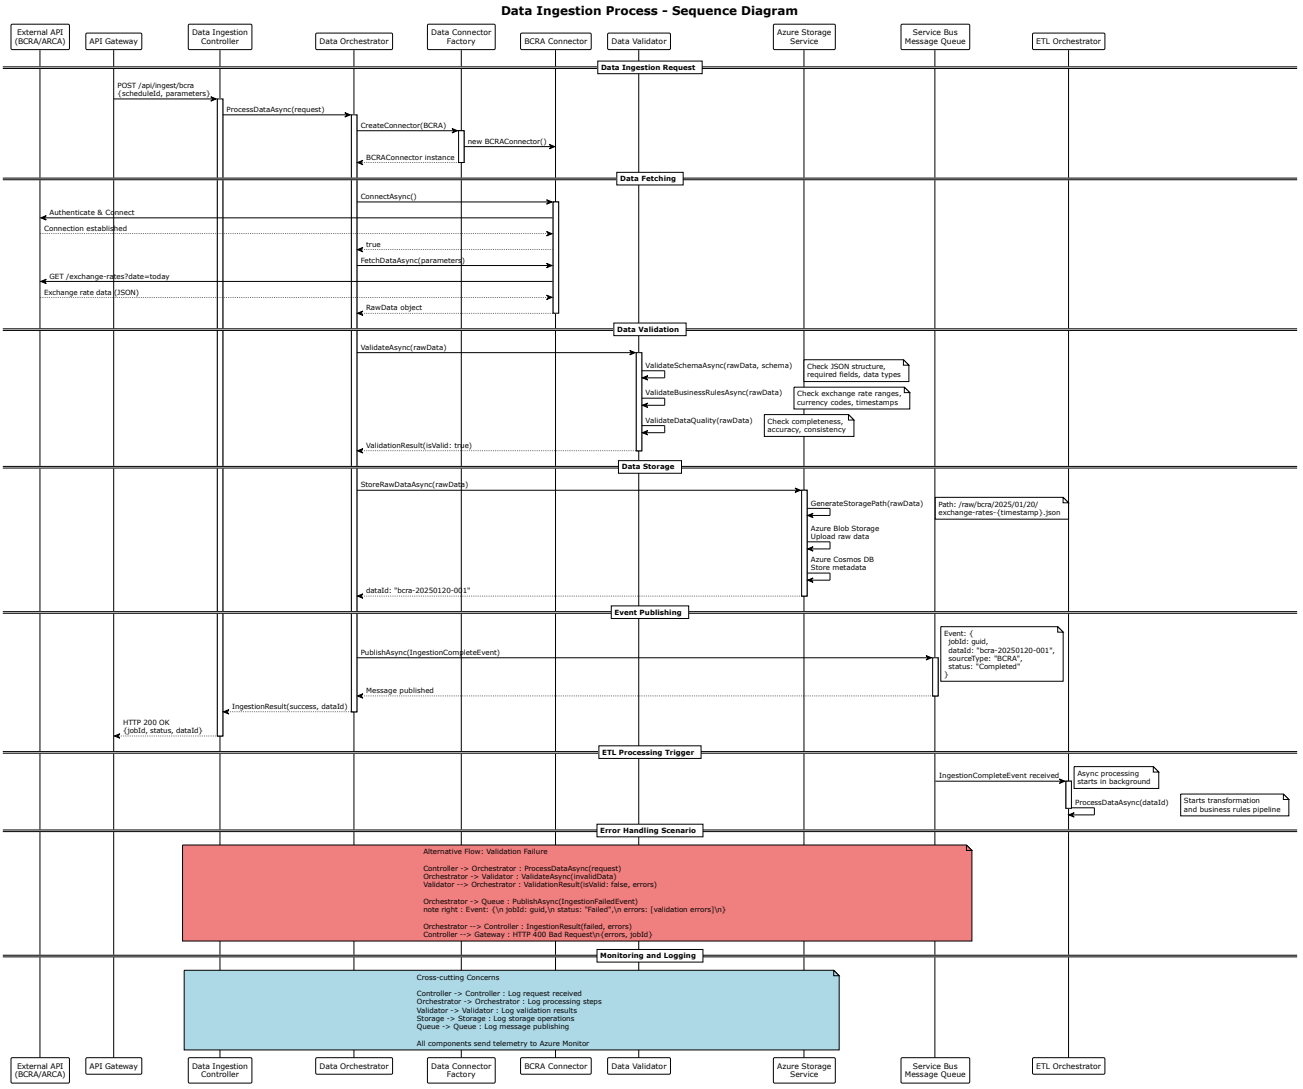

Data flow diagram

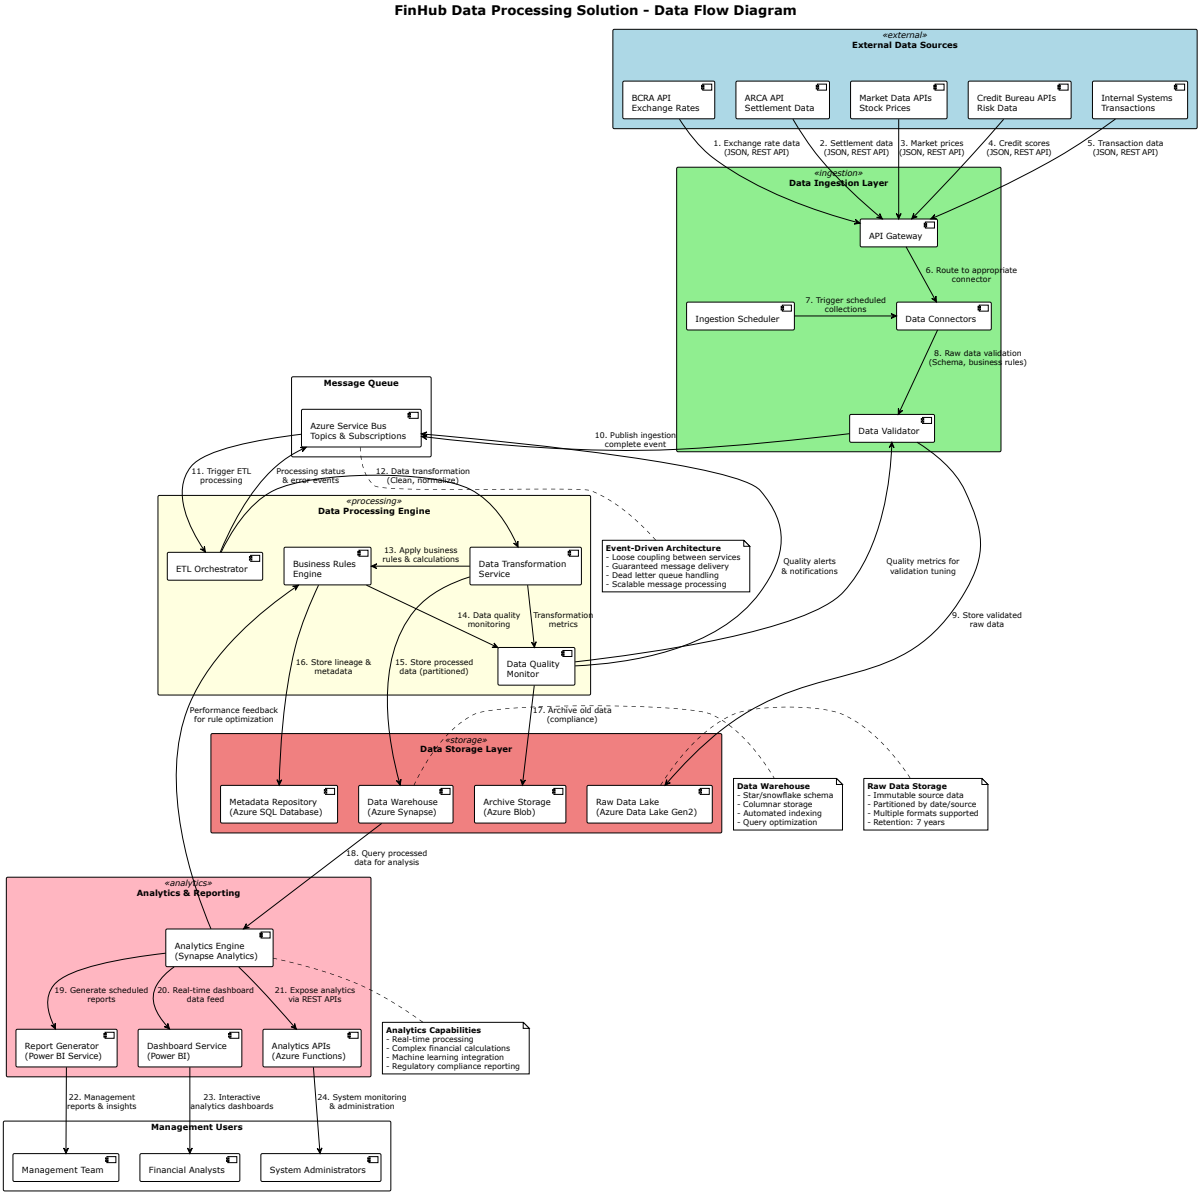

System Architecture Design Approach and Decisions

Design Summary

I built both solutions to be cloud-first, scalable, and secure—leaning heavily on Microsoft Azure's native tools. For the document classifier, I used C4 modeling to break down the architecture. For the FinHub pipeline, I went with a clean, layered approach powered by microservices and event-driven logic.

Architectural Highlights

- **Cloud-Native by Default**: Azure Functions and Synapse reduce overhead and handle scale with ease.
- **Event-Driven Messaging**: Service Bus keeps components loosely coupled and responsive.
- **Built-In Security**: Active Directory, Key Vault, and API Management ensure protection and compliance.

Technical Decisions

I went with C# for strong tooling and Azure integration. Architecture follows microservices with patterns like CQRS and event sourcing. AI features combine Azure Cognitive Services with custom models for flexibility and fast delivery.